

# Machine Learning Final Project

作者: 吳孟維 0816120

Github Link: [https://github.com/wumengwei0213/CS\\_CS20024\\_final\\_project](https://github.com/wumengwei0213/CS_CS20024_final_project)

Reference:

<https://www.kaggle.com/competitions/tabular-playground-series-aug-2022/discussion/349385>

<https://www.kaggle.com/code/johnybhiduri/tabular-playground-series-keras>

## Brief Intro.

After researching some discussion on the Kaggle, I have brief concept of this competition. I decide to design the NN model by myself in order to treat this project as a self-training and strengthen my skill in machine learning region. With looking into the data, I found that feature "loading" has a great relationship with the goal. And the features that is categorical are not much associate with the goal, so I decided to drop these feature to focus on the numeral features. The crux that I can pass the baseline is adding "model assembling" at the end of training model. I chose models that got top-5 score in the Kaggle private leaderboard and re-made the prediction.

## Model Architecture

Model: "model\_12"

Layer (type)	Output Shape	Param #
input_13 (InputLayer)	[(None, 21)]	0
normalization_12 (Normaliza tion)	(None, 21)	43
dense_60 (Dense)	(None, 256)	5632
dropout_36 (Dropout)	(None, 256)	0
batch_normalization_36 (Bat chNormalization)	(None, 256)	1024
dense_61 (Dense)	(None, 256)	65792
dropout_37 (Dropout)	(None, 256)	0
batch_normalization_37 (Bat chNormalization)	(None, 256)	1024
dense_62 (Dense)	(None, 256)	65792
dropout_38 (Dropout)	(None, 256)	0
batch_normalization_38 (Bat chNormalization)	(None, 256)	1024
dense_63 (Dense)	(None, 64)	16448
dense_64 (Dense)	(None, 1)	65
=====		
Total params: 156,844		
Trainable params: 155,265		
Non-trainable params: 1,579		

## Hyperparameters

```
optimizer = Adam with learning_rate=0.001
loss = 'binary_crossentropy'
validation_split = 0.3
batch_size = 128
epochs = 10
```

## Summary

I got 0.59022 in the end.

Name	Private Score	Public Score
inference.csv	0.59022	0.5803

Playground Prediction Competition

Tabular Playground Series - Aug 2022

Practice your ML skills on this approachable dataset!

Kaggle

1,888 teams · 4 months ago

Overview

Data

Code

Discussion

Leaderboard

Rules

Team

Submissions



Late Submission

## Submissions

You selected 0 of 2 submissions to be evaluated for your final leaderboard score. Since you selected less than 2 submission, Kaggle auto-selected up to 2 submissions from among your public best-scoring unselected submissions for evaluation. The evaluated submission with the best Private Score is used for your final score.

0/2

Submissions evaluated for final score

All		Successful	Selected	Errors	Recent ▾		
Submission and Description		Private Score ⓘ	Public Score ⓘ	Selected			
<div></div>	<div>inference.csv</div> <div>Complete (after deadline) · 2d ago</div>	0.59022	0.5803	<input type="checkbox"/>			
<div></div>	<div>inference.csv</div> <div>Complete (after deadline) · 2d ago</div>	0.58946	0.58065	<input type="checkbox"/>			

## Methodology

### Import Module

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import RobustScaler, StandardScaler
from sklearn.model_selection import KFold
# from imblearn.over_sampling import RandomOverSampler
import tensorflow as tf
from sklearn.metrics import confusion_matrix, classification_report
import warnings
warnings.filterwarnings('ignore')
```

### Data Loading

According to the research, adding columns that corresponds to if there is NA value in the features "measurement 3-6" can improve AUC value.

```
In [2]: train = pd.read_csv('train.csv')
pd.set_option('display.max_columns',None)

train['miss3'] = train['measurement_3'].isna().astype(int)*32
train['miss4'] = train['measurement_4'].isna().astype(int)*32
train['miss5'] = train['measurement_5'].isna().astype(int)*32
train['miss6'] = train['measurement_6'].isna().astype(int)*32
train.head()
```

```
Out[2]:
```

	id	product_code	loading	attribute_0	attribute_1	attribute_2	attribute_3	measurement_0	measurement_1	measurement_2	measurement_3	measur
0	0	A	80.10	material_7	material_8	9	5	7	8	4	18.040	·
1	1	A	84.89	material_7	material_8	9	5	14	3	3	18.213	·
2	2	A	82.43	material_7	material_8	9	5	12	1	5	18.057	·
3	3	A	101.07	material_7	material_8	9	5	13	2	6	17.295	·
4	4	A	188.06	material_7	material_8	9	5	9	2	8	19.346	·

### Data Preprocessing

- Pick the beneficial features for training
- Replace NA with mean value of the feature

```
In [3]: checkColumns = ['loading', 'measurement_0', 'measurement_1', 'measurement_2', 'measurement_3', 'measurement_4', 'measurement_5', 'measurme

train_df = train.copy()
for i in checkColumns :
    m = train_df[i].mean()
    train_df[i] = train_df[i].fillna(m)
```

```
Y = train_df['failure']
X = train_df[checkColumns]
```

## Deprecated Method

Add RandomOverSampler to make the data more balance

## Result

It did not enhance the score, so I remove it.

```
x_train,x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.25, shuffle = True, random_state = 144)
num_samples = int(y_train.value_counts().mean())
majority_ind = y_train[y_train == 0.0].index
samples_to_drop = y_train[majority_ind].sample(num_samples, random_state = 1).index
x_train = x_train.drop(samples_to_drop, axis = 0)
y_train = y_train.drop(samples_to_drop, axis = 0)
over_sampler = RandomOverSampler(random_state = 1)
x_train,y_train = over_sampler.fit_resample(x_train, y_train)
scaler = StandardScaler()
scaler.fit(x_train)
x_train = pd.DataFrame(scaler.transform(x_train), columns = x_train.columns, index = x_train.index)
x_test = pd.DataFrame(scaler.transform(x_test), columns = x_test.columns, index = x_test.index)
```

## Train Test Split

```
In [4]: x_train,x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.25, shuffle = True, random_state = 144)
x_train
```

```
Out[4]:
```

	loading	measurement_0	measurement_1	measurement_2	measurement_3	measurement_4	measurement_5	measurement_6	measurement_7	me
22216	109.270000	1	18	1	19.247	11.069000	17.675	17.781	10.855000	
18029	160.020000	3	7	13	18.546	10.638000	16.686	17.536	12.078000	
17046	162.310000	8	13	12	18.851	12.204000	16.031	16.599	12.715000	
3397	118.510000	11	2	2	15.033	11.413000	16.297	17.797	10.999000	
16532	147.620000	4	10	9	18.490	11.731988	20.569	17.462	13.069000	
...	...	...	...	...	...	...	...	...	...	...
17830	86.650000	5	11	10	16.651	13.058000	18.441	18.240	11.997000	
7384	109.450000	7	13	5	18.227	11.841000	17.107	18.150	11.910000	
22394	127.826233	6	11	5	18.029	9.568000	17.183	16.258	12.597000	
1468	88.290000	9	5	3	18.092	12.160000	18.499	16.669	10.590000	
25959	123.340000	9	7	5	18.669	10.919000	17.101	17.604	11.716624	

19927 rows x 21 columns

## Prediction Data Generating and Model Saving

```
In [5]: def outp(auc) :
test = pd.read_csv('test.csv')
pd_id = test['id']

test['miss3'] = test['measurement_3'].isna().astype(int)*32
test['miss4'] = test['measurement_4'].isna().astype(int)*32
test['miss5'] = test['measurement_5'].isna().astype(int)*32
test['miss6'] = test['measurement_6'].isna().astype(int)*32

for i in checkColumns :
    m = test[i].mean()
    test[i] = test[i].fillna(m)
test = test[checkColumns]

predictions = model.predict(test)
predictions = pd.DataFrame(predictions, columns=['failure'])
out = pd.concat([pd_id, predictions],axis=1)
out.to_csv(f'save/out{auc}.csv',index=False)
model.save(f'save/model_{auc}.h5')
```

## Building Model and Training

```
In [6]: while 1 :

inputs = tf.keras.Input(shape = (x_train.shape[1],))
x = tf.keras.layers.Normalization(axis=-1)(inputs)
x = tf.keras.layers.Dense(256, activation = 'relu')(x)
x = tf.keras.layers.Dropout(0.1)(x)
x = tf.keras.layers.BatchNormalization()(x)
x = tf.keras.layers.Dense(256, activation = 'relu')(x)
x = tf.keras.layers.Dropout(0.2)(x)
```

```

x = tf.keras.layers.BatchNormalization()(x)
x = tf.keras.layers.Dense(256, activation = 'relu')(x)
x = tf.keras.layers.Dropout(0.2)(x)
x = tf.keras.layers.BatchNormalization()(x)
x = tf.keras.layers.Dense(64, activation = 'relu')(x)
outputs = tf.keras.layers.Dense(1, activation = 'sigmoid')(x)

model = tf.keras.Model(inputs = inputs, outputs = outputs)
model.compile(
    optimizer = tf.keras.optimizers.Adam(learning_rate=0.001),
    loss = 'binary_crossentropy',
    metrics = [
        'accuracy',
        tf.keras.metrics.AUC(name = 'auc')
    ]
)
history = model.fit(
    x_train,
    y_train,
    validation_split = 0.3,
    batch_size = 128,
    epochs = 10,
)

print(history.history['val_auc'])

if(history.history['val_auc'][-1] > 0.58) :
    outp(str(round(history.history['val_auc'][-1],4)))
    break

```

```

Epoch 1/10
109/109 [=====] - 8s 11ms/step - loss: 0.5679 - accuracy: 0.7529 - auc: 0.5240 - val_loss: 0.5555 - val_ac
uracy: 0.7869 - val_auc: 0.5746
Epoch 2/10
109/109 [=====] - 1s 7ms/step - loss: 0.5383 - accuracy: 0.7772 - auc: 0.5376 - val_loss: 0.5141 - val_acc
uracy: 0.7919 - val_auc: 0.5639
Epoch 3/10
109/109 [=====] - 1s 7ms/step - loss: 0.5282 - accuracy: 0.7828 - auc: 0.5483 - val_loss: 0.5145 - val_acc
uracy: 0.7919 - val_auc: 0.5683
Epoch 4/10
109/109 [=====] - 1s 7ms/step - loss: 0.5276 - accuracy: 0.7841 - auc: 0.5490 - val_loss: 0.5063 - val_acc
uracy: 0.7918 - val_auc: 0.5727
Epoch 5/10
109/109 [=====] - 1s 7ms/step - loss: 0.5244 - accuracy: 0.7833 - auc: 0.5558 - val_loss: 0.5064 - val_acc
uracy: 0.7918 - val_auc: 0.5813
Epoch 6/10
109/109 [=====] - 1s 8ms/step - loss: 0.5235 - accuracy: 0.7843 - auc: 0.5522 - val_loss: 0.5064 - val_acc
uracy: 0.7911 - val_auc: 0.5712
Epoch 7/10
109/109 [=====] - 1s 7ms/step - loss: 0.5204 - accuracy: 0.7842 - auc: 0.5590 - val_loss: 0.5053 - val_acc
uracy: 0.7903 - val_auc: 0.5713
Epoch 8/10
109/109 [=====] - 1s 7ms/step - loss: 0.5201 - accuracy: 0.7843 - auc: 0.5637 - val_loss: 0.5051 - val_acc
uracy: 0.7916 - val_auc: 0.5789
Epoch 9/10
109/109 [=====] - 1s 7ms/step - loss: 0.5203 - accuracy: 0.7848 - auc: 0.5556 - val_loss: 0.5045 - val_acc
uracy: 0.7919 - val_auc: 0.5778
Epoch 10/10
109/109 [=====] - 1s 7ms/step - loss: 0.5177 - accuracy: 0.7856 - auc: 0.5662 - val_loss: 0.5076 - val_acc
uracy: 0.7916 - val_auc: 0.5774
[0.5746009349822998, 0.5639437437057495, 0.5682618021965027, 0.572697103023529, 0.581304132938385, 0.5712159872055054, 0.5713181495
666504, 0.5789045095443726, 0.5777824521064758, 0.5773547887802124]
Epoch 1/10
109/109 [=====] - 2s 8ms/step - loss: 0.5611 - accuracy: 0.7604 - auc: 0.5366 - val_loss: 0.5690 - val_acc
uracy: 0.7918 - val_auc: 0.4577
Epoch 2/10
109/109 [=====] - 1s 7ms/step - loss: 0.5403 - accuracy: 0.7793 - auc: 0.5245 - val_loss: 0.5094 - val_acc
uracy: 0.7918 - val_auc: 0.5814
Epoch 3/10
109/109 [=====] - 1s 7ms/step - loss: 0.5278 - accuracy: 0.7823 - auc: 0.5501 - val_loss: 0.5049 - val_acc
uracy: 0.7916 - val_auc: 0.5759
Epoch 4/10
109/109 [=====] - 1s 6ms/step - loss: 0.5260 - accuracy: 0.7832 - auc: 0.5535 - val_loss: 0.5065 - val_acc
uracy: 0.7918 - val_auc: 0.5638
Epoch 5/10
109/109 [=====] - 1s 7ms/step - loss: 0.5241 - accuracy: 0.7829 - auc: 0.5514 - val_loss: 0.5068 - val_acc
uracy: 0.7894 - val_auc: 0.5804
Epoch 6/10
109/109 [=====] - 1s 6ms/step - loss: 0.5237 - accuracy: 0.7839 - auc: 0.5519 - val_loss: 0.5052 - val_acc
uracy: 0.7919 - val_auc: 0.5716
Epoch 7/10
109/109 [=====] - 1s 6ms/step - loss: 0.5217 - accuracy: 0.7851 - auc: 0.5562 - val_loss: 0.5132 - val_acc
uracy: 0.7918 - val_auc: 0.5432
Epoch 8/10
109/109 [=====] - 1s 7ms/step - loss: 0.5198 - accuracy: 0.7851 - auc: 0.5575 - val_loss: 0.5076 - val_acc
uracy: 0.7919 - val_auc: 0.5736
Epoch 9/10
109/109 [=====] - 1s 7ms/step - loss: 0.5199 - accuracy: 0.7851 - auc: 0.5703 - val_loss: 0.5063 - val_acc
uracy: 0.7918 - val_auc: 0.5724
Epoch 10/10
109/109 [=====] - 1s 7ms/step - loss: 0.5192 - accuracy: 0.7851 - auc: 0.5655 - val_loss: 0.5058 - val_acc
uracy: 0.7918 - val_auc: 0.5809
[0.45770102739334106, 0.5813931822776794, 0.575866580094604, 0.5637527108192444, 0.5804040431976318, 0.5715550780296326, 0.5431562
066078186, 0.5736348032951355, 0.5723981857299805, 0.5808916687965393]
650/650 [=====] - 1s 1ms/step

```

```
In [7]: model.summary()
```

Model: "model\_1"

Layer (type)	Output Shape	Param #
=====		
input_2 (InputLayer)	[(None, 21)]	0
normalization_1 (Normalizat ion)	(None, 21)	43
dense_5 (Dense)	(None, 256)	5632
dropout_3 (Dropout)	(None, 256)	0
batch_normalization_3 (Batc hNormalization)	(None, 256)	1024
dense_6 (Dense)	(None, 256)	65792
dropout_4 (Dropout)	(None, 256)	0
batch_normalization_4 (Batc hNormalization)	(None, 256)	1024
dense_7 (Dense)	(None, 256)	65792
dropout_5 (Dropout)	(None, 256)	0
batch_normalization_5 (Batc hNormalization)	(None, 256)	1024
dense_8 (Dense)	(None, 64)	16448
dense_9 (Dense)	(None, 1)	65
=====		
Total params: 156,844		
Trainable params: 155,265		
Non-trainable params: 1,579		

## Model Evaluating

```
In [8]: results = model.evaluate(x_test, y_test, verbose = 0)
print('Test Loss : {:.4f}%'.format(results[0]*100))
print('Test Accuracy : {:.3f}%'.format(results[1]*100))
print('Test AUC : {:.4f}%'.format(results[2]*100))
```

Test Loss : 50.8950%  
Test Accuracy : 78.729%  
Test AUC : 59.6999%

```
In [9]: y_pred = np.array(model.predict(x_test) >= 0.5, dtype = np.int)
```

208/208 [=====] - 0s 1ms/step

```
In [10]: cm = confusion_matrix(y_test, y_pred)
clr = classification_report(y_test, y_pred, target_names = ['Failure', 'Success'])
sns.heatmap(cm, annot = True, fmt = 'g', cbar = False, cmap = 'Blues')
plt.xticks(ticks = (0.5, 1.5), labels = ['Failure', 'Success'])
plt.yticks(ticks = (0.5, 1.5), labels = ['Failure', 'Success'])
plt.show()
print(f'Classification Report : \n{clr}')
```



```
Classification Report :
      precision    recall  f1-score   support

   Failure       0.79      1.00      0.88       5230
   Success       0.00      0.00      0.00       1413

   accuracy          0.79      0.79      0.79       6643
  macro avg       0.39      0.50      0.44       6643
 weighted avg       0.62      0.79      0.69       6643
```

## Model Assembling

- Choose models that got top-5 score in the Kaggle private leaderboard and re-made the prediction.

```
pre = []
test = pd.read_csv('test.csv')
pd_id = test['id']
dirlist = os.listdir('infer')
for doc in dirlist:
    df = pd.read_csv('infer/'+doc)
    # Get the value after ranking with the failure rate
    rank = df['failure'].rank(axis=0).rename(doc)
    pd_id = pd.concat([pd_id, rank],axis=1)

pd_id['min'] = pd_id.drop('id', axis=1).min(axis=1)
pd_id['max'] = pd_id.drop('id', axis=1).max(axis=1)
pd_id['mean'] = pd_id.drop('id', axis=1).mean(axis=1)
pd_id['median'] = pd_id.drop('id', axis=1).median(axis=1)

df = pd_id['max'].copy()
# Normalization
normalized_df=(df-df.min())/(df.max()-df.min())
out = pd.concat([test['id'], normalized_df.rename('failure')],axis=1)

# Generate prediction data
out.to_csv('inference.csv',index=False)
```

- In the end, I chose the max ranking value to do the prediction in those prediction data after several experiments.

	id	model_5847.csv	model_5856.csv	model_5859.csv	model_5869.csv	model_587.csv	model_5872.csv	model_5875.csv	min	max	mean	median
0	26570	9771.0	10200.0	9965.0	10108.0	9888.0	10051.0	10080.0	9771.0	10200.0	10003.777778	10027.388889
1	26571	6448.0	7112.0	6346.0	6425.0	7290.0	6438.0	6714.0	6346.0	7290.0	6712.111111	6580.055556
2	26572	6928.0	6926.0	6444.0	7063.0	6750.0	6765.0	6051.0	6051.0	7063.0	6671.222222	6757.500000
3	26573	9539.0	9117.0	9494.0	9082.0	8977.0	9632.0	9979.0	8977.0	9979.0	9419.555556	9456.777778
4	26574	20248.0	19939.0	19983.0	19941.0	20320.5	20005.0	20008.0	19939.0	20320.5	20078.222222	20006.500000
...	...	...	...	...	...	...	...	...	...	...	...	...
20770	47340	14876.0	14977.0	15070.0	15003.0	14360.0	15148.0	13544.0	13544.0	15148.0	14630.000000	14926.500000
20771	47341	1402.0	1239.0	1582.0	1240.0	1292.0	1352.0	1509.0	1239.0	1582.0	1381.888889	1366.944444
20772	47342	3227.0	1674.0	4167.0	3181.0	1298.0	917.0	1436.0	917.0	4167.0	2331.555556	2002.777778
20773	47343	11604.0	11322.0	11379.0	11835.0	11405.0	11882.0	12043.0	11322.0	12043.0	11648.333333	11626.166667
20774	47344	2698.0	2681.0	2912.0	2885.0	2502.0	2819.0	2493.0	2493.0	2912.0	2710.555556	2704.277778

20775 rows × 12 columns